

# **CS:APP Chapter 4**

## **Computer Architecture**

# **Overview**

**Yuan Tang**

***Adapted from CMU course 15-213***

# Class Staff

Instructor: 唐渊

Email: [yuantang@fudan.edu.cn](mailto:yuantang@fudan.edu.cn)

Office hour: by appointment

TA: 陆溢超, 欧承祖

Email: [\[luyc13, oucz14\]@fudan.edu.cn](mailto:[luyc13, oucz14]@fudan.edu.cn)

Office hour: by appointment

# Grading

## Exams(60%)

- Mid term (30%)
- Final (30%)
- All exams are open books/open notes.

## Find bugs

- Including online pptx, textbook, online code
- 10 points each, 5 points if doesn't find the exact reason
- Only credit the first people finding the bug

## Using MOOC++

- Including raising / answering questions, suggestions to MOOC++
- 2 points for each good question (judged by TAs or instructor), 2 points for a totally matched answer (to normalize), 2 points for a good suggestion to MOOC++

# Course Outline

## Background

- Instruction sets
- Logic design

## Sequential Implementation

- A simple, but not very fast processor design

## Pipelining

- Get more things running simultaneously

## Pipelined Implementation

- Make it work

## Advanced Topics

- Performance analysis
- High performance processor design

# Coverage

## Our Approach

- **Work through designs for particular instruction set**
  - Y86-64 – a simplified version of the Intel x86-64
  - If you know one, you more-or-less know them all
- **Work at “microarchitectural” level**
  - **Assemble basic hardware blocks into overall processor structure**
    - » Memories, functional units, etc.
  - **Surround by control logic to make sure each instruction flows through properly**

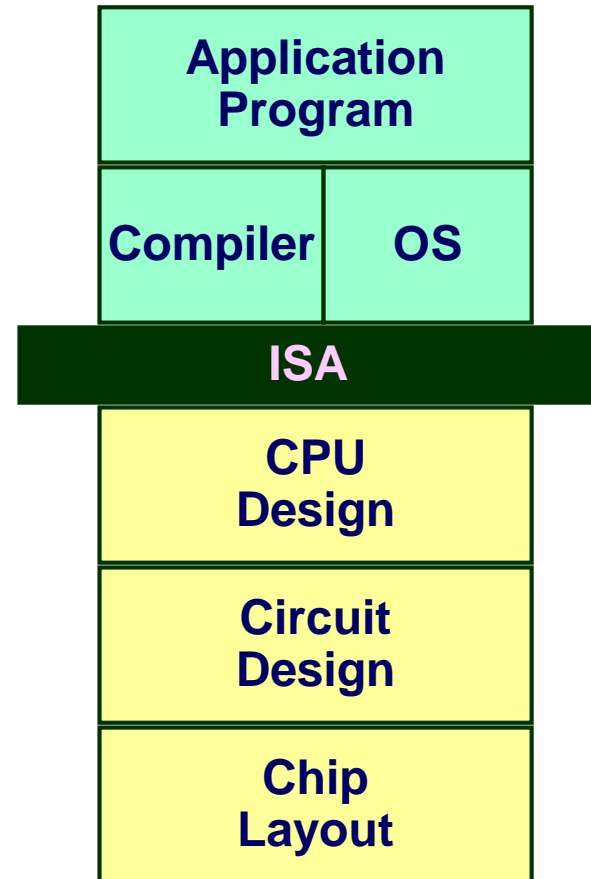
# Instruction Set Architecture

## Assembly Language View

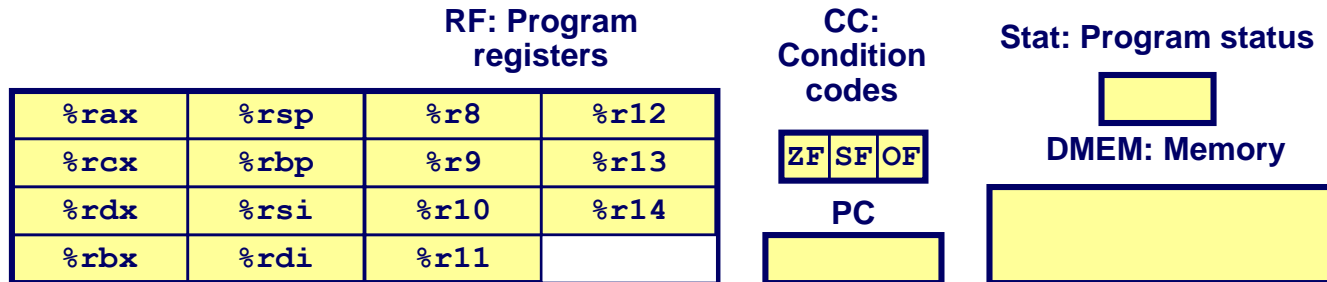
- **Processor state**
  - Registers, memory, ...
- **Instructions**
  - `addq, pushq, ret, ...`
  - How instructions are encoded as bytes

## Layer of Abstraction

- **Above: how to program machine**
  - Processor executes instructions in a sequence
- **Below: what needs to be built**
  - Use variety of tricks to make it run fast
  - E.g., execute multiple instructions simultaneously



# Y86-64 Processor State



## ■ Program Registers

- 15 registers (omit %r15). Each 64 bits

## ■ Condition Codes

- Single-bit flags set by arithmetic or logical instructions

» ZF: Zero                      SF: Negative                      OF: Overflow

## ■ Program Counter

- Indicates address of next instruction

## ■ Program Status

- Indicates either normal operation or some error condition

## ■ Memory

- Byte-addressable storage array
- Words stored in little-endian byte order

# Y86-64 Instruction Set #1

Byte	0	1	2	3	4	5	6	7	8	9
halt	0	0								
nop	1	0								
cmovXX rA, rB	2	fn	rA	rB						
irmovq V, rB	3	0	F	rB	V					
rmmovq rA, D(rB)	4	0	rA	rB	D					
mrmmovq D(rB), rA	5	0	rA	rB	D					
OPq rA, rB	6	fn	rA	rB						
jXX Dest	7	fn	Dest							
call Dest	8	0	Dest							
ret	9	0								
pushq rA	A	0	rA	F						
popq rA	B	0	rA	F						



# Y86-64 Instructions

## Format

- 1–10 bytes of information read from memory
  - Can determine instruction length from first byte
  - Not as many instruction types, and simpler encoding than with x86-64
- Each accesses and modifies some part(s) of the program state

# Y86-64 Instruction Set #2

Byte

	0	1	2	3	4	5	6	
halt	0	0						
nop	1	0						
cmovXX rA, rB	2	fn	rA	rB				rrmovq 2 0
irmovq V, rB	3	0	F	rB				cmovle 2 1
rrmovq rA, D(rB)	4	0	rA	rB				cmovl 2 2
mrmovq D(rB), rA	5	0	rA	rB				cmove 2 3
OPq rA, rB	6	fn	rA	rB				cmovne 2 4
jXX Dest	7	fn						cmovge 2 5
call Dest	8	0						cmovg 2 6
ret	9	0						
pushq rA	A	0	rA	F				
popq rA	B	0	rA	F				

# Y86-64 Instruction Set #3

Byte	0	1	2	3	4	5	6	7	8	9				
halt	0	0												
nop	1	0												
cmovXX rA, rB	2	fn	rA	rB										
irmovq V, rB	3	0	F	rB	V									
rmmovq rA, D(rB)	4	0	rA	rB	D									
mrmmovq D(rB), rA	5	0	rA	rB	D									
OPq rA, rB	6	fn	rA	rB	<div><div></div><div>addq</div><div>6</div><div>0</div></div> <div><div></div><div>subq</div><div>6</div><div>1</div></div> <div><div></div><div>andq</div><div>6</div><div>2</div></div> <div><div></div><div>xorq</div><div>6</div><div>3</div></div>									
jXX Dest	7	fn	Dest											
call Dest	8	0	Dest											
ret	9	0												
pushq rA	A	0	rA	F										
popq rA	B	0	rA	F										

CS:A

# Y86-64 Instruction Set #4

Byte	0	1	2	3	4	5	6	7		
halt	0	0								
nop	1	0								
cmovXX rA, rB	2	fn	rA	rB						
irmovq V, rB	3	0	F	rB				V		
rmmovq rA, D(rB)	4	0	rA	rB				D		
mrmovq D(rB), rA	5	0	rA	rB				D		
OPq rA, rB	6	fn	rA	rB						
jXX Dest	7	fn								
call Dest	8	0								
ret	9	0								
pushq rA	A	0	rA	F						
popq rA	B	0	rA	F						

jmp	7	0
jle	7	1
j1	7	2
je	7	3
jne	7	4
jge	7	5
jg	7	6

# Encoding Registers

Each register has 4-bit ID

%rax	0	%r8	8
%rcx	1	%r9	9
%rdx	2	%r10	A
%rbx	3	%r11	B
%rsp	4	%r12	C
%rbp	5	%r13	D
%rsi	6	%r14	E
%rdi	7	No Register	F

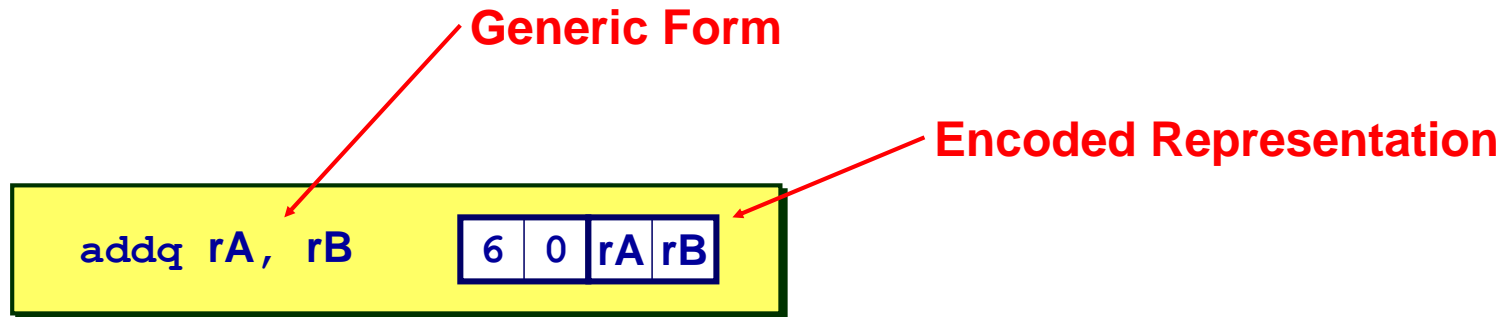
- Same encoding as in x86-64

**Register ID 15 (0xF) indicates “no register”**

- Will use this in our hardware design in multiple places

# Instruction Example

## Addition Instruction



- Add value in register rA to that in register rB
  - Store result in register rB
  - Note that Y86-64 only allows addition to be applied to register data
- Set condition codes based on result
- e.g., `addq %rax,%rsi` Encoding: 60 06
- Two-byte encoding
  - First indicates instruction type
  - Second gives source and destination registers

# Arithmetic and Logical Operations

Instruction Code

Function Code


Add

`addq rA, rB`




Subtract (rA from rB)

`subq rA, rB`




And

`andq rA, rB`



Exclusive-Or

`xorq rA, rB`



- Refer to generically as “OPq”
- Encodings differ only by “function code”
  - Low-order 4 bytes in first instruction word
- Set condition codes as side effect

# Move Operations

Register → Register

`rrmovq rA, rB`



Immediate → Register

`irmovq V, rB`



Register → Memory

`rmmovq rA, D(rB)`



Memory → Register

`mrmmovq D(rB), rA`



- Like the x86-64 `movq` instruction
- Simpler format for memory addresses
- Give different names to keep them distinct



# Move Instruction Examples

## X86-64

```
movq $0xabcd, %rdx
```

Encoding: 30 82 cd ab 00 00 00 00 00 00

```
movq %rsp, %rbx
```

Encoding: 20 43

```
movq -12(%rbp), %rcx
```

Encoding: 50 15 f4 ff ff ff ff ff ff ff

```
movq %rsi, 0x41c(%rsp)
```

Encoding: 40 64 1c 04 00 00 00 00 00 00

## Y86-64

```
irmovq $0xabcd, %rdx
```

```
rrmovq %rsp, %rbx
```

```
mrmovq -12(%rbp), %rcx
```

```
rmmovq %rsi, 0x41c(%rsp)
```

# Conditional Move Instructions

## Move Unconditionally

`rrmovq rA, rB`



## Move When Less or Equal

`cmovle rA, rB`



## Move When Less

`cmovl rA, rB`



## Move When Equal

`cmove rA, rB`



## Move When Not Equal

`cmovne rA, rB`




## Move When Greater or Equal

`cmovge rA, rB`



## Move When Greater

`cmovg rA, rB`



- Refer to generically as “`cmovXX`”
- Encodings differ only by “function code”
- Based on values of condition codes
- Variants of `rrmovq` instruction
  - (Conditionally) copy value from source to destination register

# Jump Instructions

## Jump (Conditionally)



- Refer to generically as “jxx”
- Encodings differ only by “function code” fn
- Based on values of condition codes
- Same as x86-64 counterparts
- Encode full destination address
  - Unlike PC-relative addressing seen in x86-64

# Jump Instructions

## Jump Unconditionally

**jmp Dest**    7   0    Dest

## Jump When Less or Equal

**jle Dest**    7   1    Dest

## Jump When Less

**jlt Dest**    7   2    Dest

## Jump When Equal

**je Dest**    7   3    Dest

## Jump When Not Equal

**jne Dest**    7   4    Dest


## Jump When Greater or Equal

**jge Dest**    7   5    Dest

## Jump When Greater

**jg Dest**    7   6    Dest

# Y86-64 Program Stack



- Region of memory holding program data
- Used in Y86-64 (and x86-64) for supporting procedure calls
- Stack top indicated by `%rsp`
  - Address of top stack element
- Stack grows toward lower addresses
  - Top element is at highest address in the stack
  - When pushing, must first decrement stack pointer
  - After popping, increment stack pointer

# Stack Operations

`pushq rA`

A	0	rA	F
---	---	----	---

- Decrement `%rsp` by 8
- Store word from `rA` to memory at `%rsp`
- Like x86-64

`popq rA`

B	0	rA	F
---	---	----	---

- Read word from memory at `%rsp`
- Save in `rA`
- Increment `%rsp` by 8
- Like x86-64

# Subroutine Call and Return

`call Dest`

8

0

Dest

- Push address of next instruction onto stack
- Start executing instructions at Dest
- Like x86-64

`ret`

9

0

- Pop value from stack
- Use as address for next instruction
- Like x86-64

# Miscellaneous Instructions

`nop`



- Don't do anything

`halt`



- Stop executing instructions
- x86-64 has comparable instruction, but can't execute it in user mode
- We will use it to stop the simulator
- Encoding ensures that program hitting memory initialized to zero will halt



# Status Conditions

Mnemonic	Code
AOK	1

- Normal operation

Mnemonic	Code
HLT	2

- Halt instruction encountered

Mnemonic	Code
ADR	3

- Bad address (either instruction or data) encountered

Mnemonic	Code
INS	4

- Invalid instruction encountered

## Desired Behavior

- If AOK, keep going
- Otherwise, stop program execution

# Writing Y86-64 Code


## Try to Use C Compiler as Much as Possible

- Write code in C
- Compile for x86-64 with `gcc -Og -S`
- Transliterate into Y86-64
- *Modern compilers make this more difficult*

## Coding Example

- Find number of elements in null-terminated list

```
int len1(int a[]);
```



# Y86-64 Code Generation Example

## First Try

- Write typical array code

```
/* Find number of elements in
   null-terminated list */
long len(long a[])
{
    long len;
    for (len = 0; a[len]; len++)
        ;
    return len;
}
```

- Compile with `gcc -Og -S`

## Problem

- Hard to do array indexing on Y86-64
  - Since don't have scaled addressing modes

L3:

```
addq $1,%rax
cmpq $0, (%rdi,%rax,8)
jne L3
```

# Y86-64 Code Generation Example #2

## Second Try

- Write C code that mimics expected Y86-64 code

```
long len2(long *a)
{
    long ip = (long) a;
    long val = *(long *) ip;
    long len = 0;
    while (val) {
        ip += sizeof(long);
        len++;
        val = *(long *) ip;
    }
    return len;
}
```

## Result

- Compiler generates exact same code as before!
- Compiler converts both versions into same intermediate form

# Y86-64 Code Generation Example #3

```
len:
    irmovq $1, %r8          # Constant 1
    irmovq $8, %r9          # Constant 8
    irmovq $0, %rax         # len = 0
    mrmovq (%rdi), %rdx     # val = *a
    andq %rdx, %rdx         # Test val
    je Done                 # If zero, goto Done

Loop:
    addq %r8, %rax          # len++
    addq %r9, %rdi          # a++
    mrmovq (%rdi), %rdx     # val = *a
    andq %rdx, %rdx         # Test val
    jne Loop                # If !0, goto Loop

Done:
    ret
```

Register	Use
%rdi	a
%rax	len
%rdx	val
%r8	1
%r9	8

# Y86-64 Sample Program Structure #1

```
init:                                # Initialization
    . . .
    call Main
    halt

    .align 8                          # Program data
array:
    . . .

Main:                                # Main function
    . . .
    call len    . . .

len:                                  # Length function
    . . .

    .pos 0x100                        # Placement of stack
Stack:
```

- Program starts at address 0
- Must set up stack
  - Where located
  - Pointer values
  - Make sure don't overwrite code!
- Must initialize data

# Y86-64 Program Structure #2

```
init:
    # Set up stack pointer
    irmovq Stack, %rsp
    # Execute main program
    call Main
    # Terminate
    halt

# Array of 4 elements + terminating 0
    .align 8
Array:
    .quad 0x000d000d000d000d
    .quad 0x00c000c000c000c0
    .quad 0x0b000b000b000b00
    .quad 0xa000a000a000a000
    .quad 0
```

- Program starts at address 0
- Must set up stack
- Must initialize data
- Can use symbolic names

# Y86-64 Program Structure #3

```
Main:
    irmovq array,%rdi
    # call len(array)
    call len
    ret
```

## Set up call to len

- Follow x86-64 procedure conventions
- Push array address as argument



# Assembling Y86-64 Program

```
unix> yas len.yas
```

- Generates “object code” file `len.yo`
  - Actually looks like disassembler output

```
0x054:          | len:
0x054: 30f8010000000000000000 |    irmovq $1, %r8          # Constant 1
0x05e: 30f9080000000000000000 |    irmovq $8, %r9          # Constant 8
0x068: 30f0000000000000000000 |    irmovq $0, %rax         # len = 0
0x072: 5027000000000000000000 |    mrmovq (%rdi), %rdx     # val = *a
0x07c: 6222          |    andq %rdx, %rdx         # Test val
0x07e: 73a0000000000000000000 |    je Done                 # If zero, goto Done
0x087:          | Loop:
0x087: 6080          |    addq %r8, %rax          # len++
0x089: 6097          |    addq %r9, %rdi          # a++
0x08b: 5027000000000000000000 |    mrmovq (%rdi), %rdx     # val = *a
0x095: 6222          |    andq %rdx, %rdx         # Test val
0x097: 7487000000000000000000 |    jne Loop                # If !0, goto Loop
0x0a0:          | Done:
0x0a0: 90            |    ret
```

# Simulating Y86-64 Program

```
unix> yis len.yo
```

## ■ Instruction set simulator

- Computes effect of each instruction on processor state
- Prints changes in state from original

```
Stopped in 33 steps at PC = 0x13.  Status 'HLT', CC Z=1 S=0 O=0
```

```
Changes to registers:
```

%rax:	0x0000000000000000	0x0000000000000004
%rsp:	0x0000000000000000	0x0000000000000100
%rdi:	0x0000000000000000	0x0000000000000038
%r8:	0x0000000000000000	0x0000000000000001
%r9:	0x0000000000000000	0x0000000000000008

```
Changes to memory:
```

0x00f0:	0x0000000000000000	0x0000000000000053
0x00f8:	0x0000000000000000	0x0000000000000013

# CISC Instruction Sets

- Complex Instruction Set Computer
- IA32 is example

## Stack-oriented instruction set

- Use stack to pass arguments, save program counter
- Explicit push and pop instructions

## Arithmetic instructions can access memory

- `addq %rax, 12(%rbx,%rcx,8)`
  - requires memory read and write
  - Complex address calculation

## Condition codes

- Set as side effect of arithmetic and logical instructions

## Philosophy

- Add instructions to perform “typical” programming tasks

# RISC Instruction Sets

- Reduced Instruction Set Computer
- Internal project at IBM, later popularized by Hennessy (Stanford) and Patterson (Berkeley)

## Fewer, simpler instructions

- Might take more to get given task done
- Can execute them with small and fast hardware

## Register-oriented instruction set

- Many more (typically 32) registers
- Use for arguments, return pointer, temporaries

## Only load and store instructions can access memory

- Similar to Y86-64 `mrmovq` and `rmmovq`

## No Condition codes

- Test instructions return 0/1 in register

# MIPS Registers

\$0	\$0	Constant 0	\$16	\$s0	Callee Save Temporaries: May not be overwritten by called procedures
\$1	\$at	Reserved Temp.	\$17	\$s1	
\$2	\$v0	Return Values	\$18	\$s2	
\$3	\$v1		\$19	\$s3	
\$4	\$a0	Procedure arguments	\$20	\$s4	
\$5	\$a1		\$21	\$s5	Caller Save Temp
\$6	\$a2		\$22	\$s6	
\$7	\$a3		\$23	\$s7	
\$8	\$t0	Caller Save Temporaries: May be overwritten by called procedures	\$24	\$t8	
\$9	\$t1		\$25	\$t9	
\$10	\$t2		\$26	\$k0	Reserved for Operating Sys
\$11	\$t3		\$27	\$k1	
\$12	\$t4		\$28	\$gp	Global Pointer
\$13	\$t5		\$29	\$sp	Stack Pointer
\$14	\$t6		\$30	\$s8	Callee Save Temp
\$15	\$t7		\$31	\$ra	Return Address

# MIPS Instruction Examples

## R-R

Op	Ra	Rb	Rd	00000	Fn
----	----	----	----	-------	----

`addu $3,$2,$1`                      # Register add:  $\$3 = \$2 + \$1$

## R-I

Op	Ra	Rb	Immediate
----	----	----	-----------

`addu $3,$2, 3145`                      # Immediate add:  $\$3 = \$2 + 3145$

`sll $3,$2,2`                              # Shift left:  $\$3 = \$2 \ll 2$

## Branch

Op	Ra	Rb	Offset
----	----	----	--------

`beq $3,$2,dest`                      # Branch when  $\$3 = \$2$

## Load/Store

Op	Ra	Rb	Offset
----	----	----	--------

`lw $3,16($2)`                              # Load Word:  $\$3 = M[\$2 + 16]$

`sw $3,16($2)`                              # Store Word:  $M[\$2 + 16] = \$3$

# CISC vs. RISC

## Original Debate

- Strong opinions!
- CISC proponents---easy for compiler, fewer code bytes
- RISC proponents---better for optimizing compilers, can make run fast with simple chip design

## Current Status

- For desktop processors, choice of ISA not a technical issue
  - With enough hardware, can make anything run fast
  - Code compatibility more important
- x86-64 adopted many RISC features
  - More registers; use them for argument passing
- For embedded processors, RISC makes sense
  - Smaller, cheaper, less power
  - Most cell phones use ARM processor

# Summary

## Y86-64 Instruction Set Architecture

- Similar state and instructions as x86-64
- Simpler encodings
- Somewhere between CISC and RISC

## How Important is ISA Design?

- Less now than before
  - With enough hardware, can make almost anything go fast